

Databases and Advanced Data Techniques

Tarapong Sreenuch

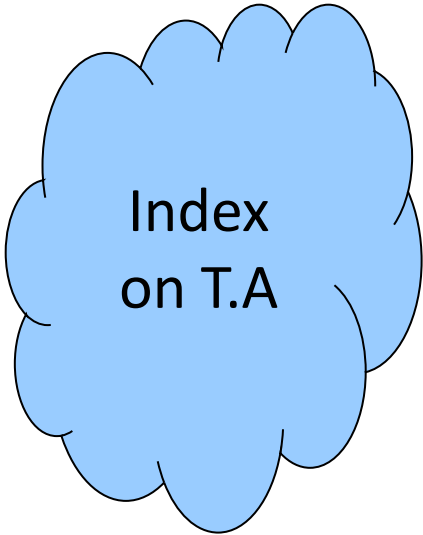
Indexes and Transactions

Indexes

- Primary mechanism to get improved performance on a database
- Persistent data structure, stored in a database
- Many interesting implementation techniques
 - In this course, we will only be focusing on user/application perspective.



Functionality



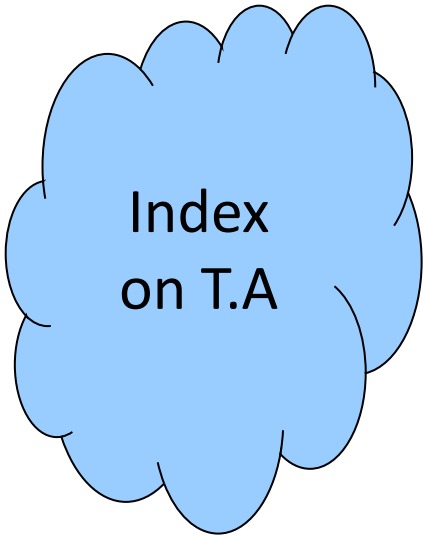
T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Functionality

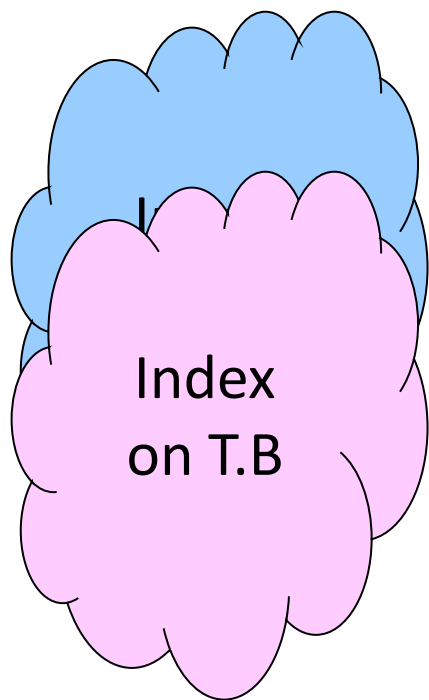
$T.A = 'cat'$



T			
	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Functionality



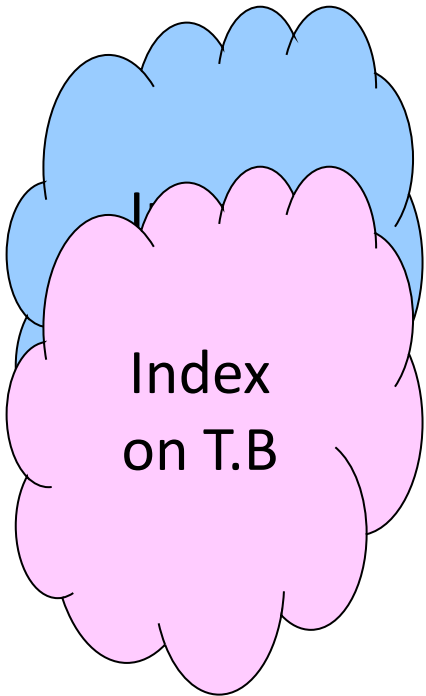
T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Functionality

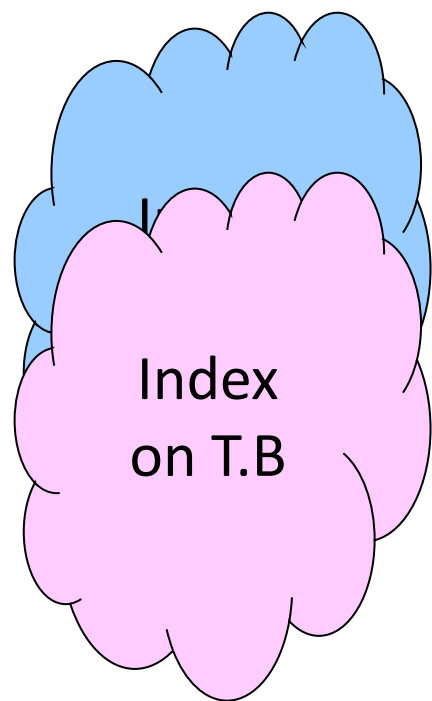
$T.B = 9$



T			
	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Functionality

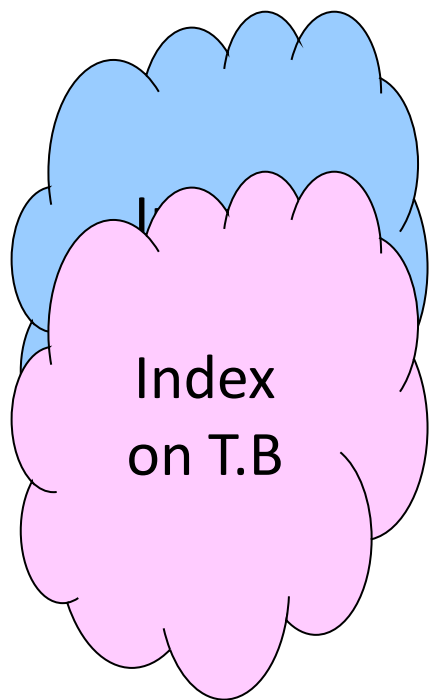


T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Functionality

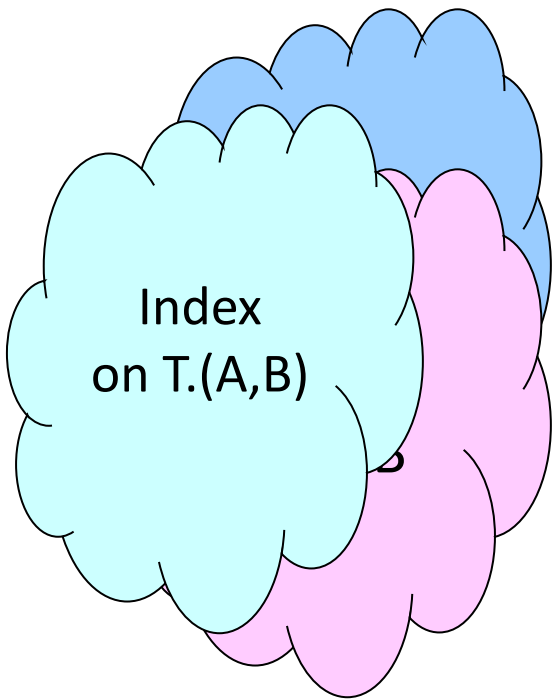


T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Functionality



$T.A = 'cat' \text{ AND } T.B = 8$

T

	A	B	C
1	cat	2	...
2	dog	5	...
3	cow	1	...
4	dog	9	...
5	cat	2	...
6	cat	8	...
7	cow	6	...



Utility

- Index = difference between full table scans and immediate location of tuples
 - Order of magnitude performance difference
- Underlying data structures
 - Balanced trees (B trees, B+ trees)
 - Hash tables



Index Operation

```
SELECT sName  
FROM Student  
WHERE sID = 18942
```

Index on sID

Many DMBS build indexes automatically on **PRIMARY KEY** (and sometimes **UNIQUE**) attributes.



Index Search Operation

```
SELECT sID  
FROM Student  
WHERE sName = 'Mary' AND GPA > 3.9
```

Index on sName \leftarrow Hash- or Tree based indexes are required in order to support equality search.

Index on GPA \leftarrow Tree based indexes are required in order to support inequality search.

Index on (sName, GPA)



Index Search Operation

```
SELECT sName, cName  
FROM Student, Apply  
WHERE Student.sID = Apply.sID
```

Index *Index*

Query Planning & Optimisation



Downsides of Indexes

Benefit of an index depends on:

- Extra Space \leftarrow *Marginal Downside*
- Index creation \leftarrow *Medium Downside*
- Index Maintenance \leftarrow *Cost Benefit Tradeoffs*



Which Indexes To Create?

Benefit of an index depends on:

- Size of table (and possibly layout)
- Data distributions
- Query vs. update load

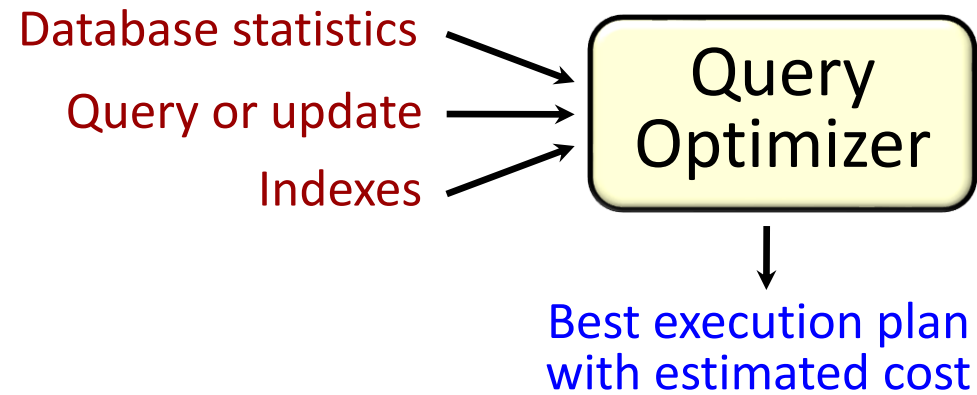


Physical Design Advisors

Input: Database (Statistics) and Workload

Output: Recommendation Indexes

We're looking for Indexes that their benefits outweigh the drawbacks.



SQL Syntax

```
CREATE INDEX IndexName ON T(A)
```

```
CREATE INDEX IndexName ON T(A1,A2,...,An)
```

```
CREATE UNIQUE INDEX IndexName ON T(A)
```

```
DROP INDEX IndexName
```



Recap: Indexes

- Primary mechanism to get improved performance on a database
- Persistent data structure, stored in database
- Ways to implementation it:
 - Hash, B-Tree, etc.



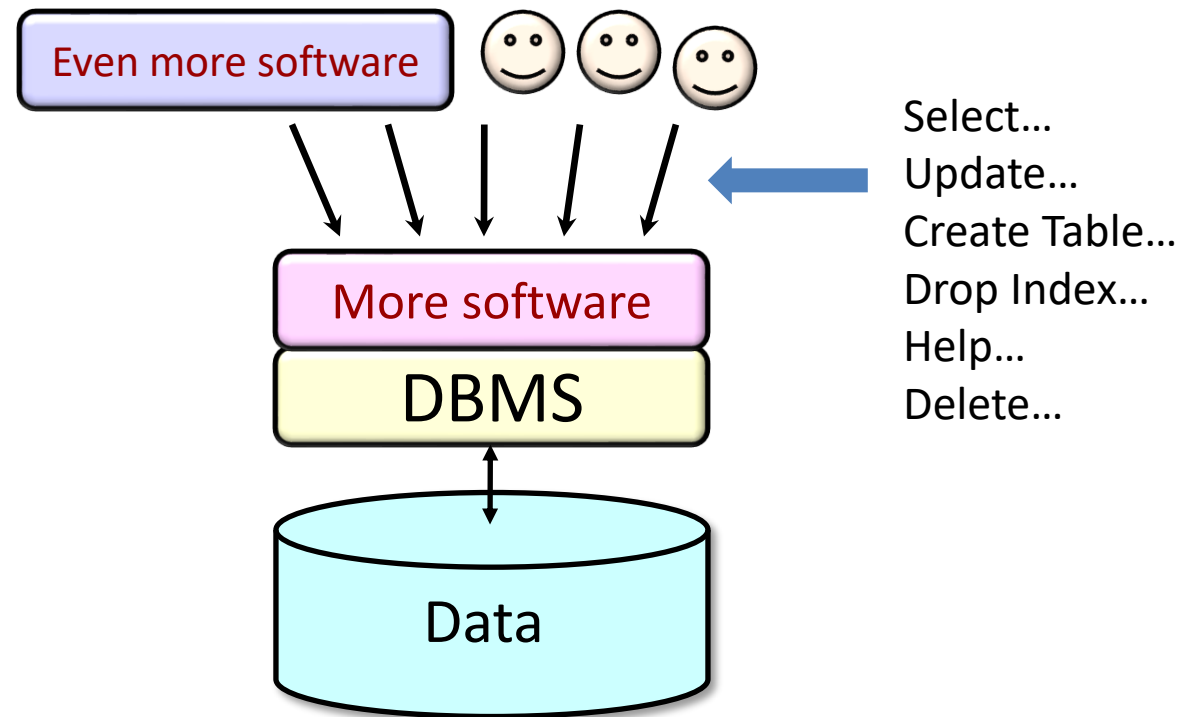
Transactions

Motivated by two independent requirements

- Concurrent database access
- Resilience to system failures



Concurrent Database Access



Concurrent Access: Attribute-Level Inconsistency

```
UPDATE College SET enrollment = enrollment + 1000  
WHERE cName = 'University of London'
```

... concurrent with ...

```
UPDATE College SET enrollment = enrollment + 1500  
WHERE cName = 'University of London'
```



Concurrent Access: Tuple-Level Inconsistency

```
UPDATE Apply SET major = 'CS'  
WHERE sID = 123
```

... concurrent with ...

```
UPDATE Apply SET decision = 'Y'  
WHERE sID = 123
```



Concurrent Access: Table-Level Inconsistency

```
UPDATE Apply SET decision = 'Y'  
WHERE sID IN ( SELECT sID  
                FROM Student  
                WHERE GPA > 3.9 )
```

... concurrent with ...

```
UPDATE Student SET GPA = (1.1)*GPA  
WHERE sizeHS > 2500
```



Concurrent Access: Multi-Statement Inconsistency

```
INSERT INTO Archive  
  SELECT * FROM Apply WHERE decision = 'N';  
DELETE FROM Apply WHERE decision = 'N';
```

... concurrent with ...

```
SELECT COUNT(*) FROM Apply;  
SELECT COUNT(*) FROM Archive;
```



Concurrent Goal

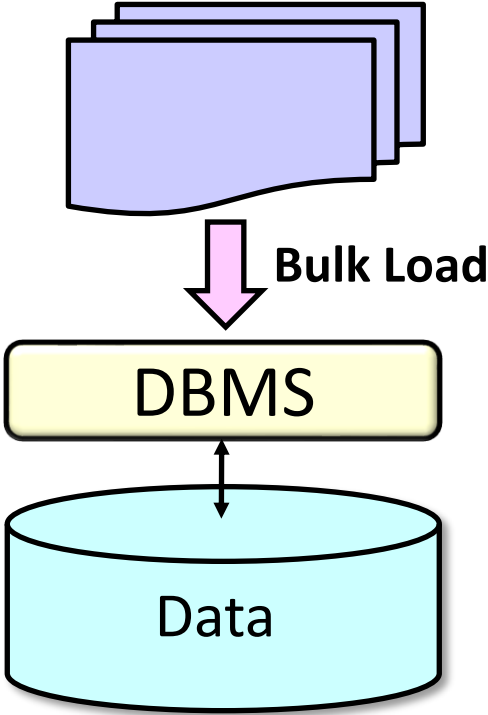
Execute sequence of SQL statements so they appear to be running in isolation

- A simple solution is to execute them in isolation.

But still we want to enable concurrency whenever safe to do so.

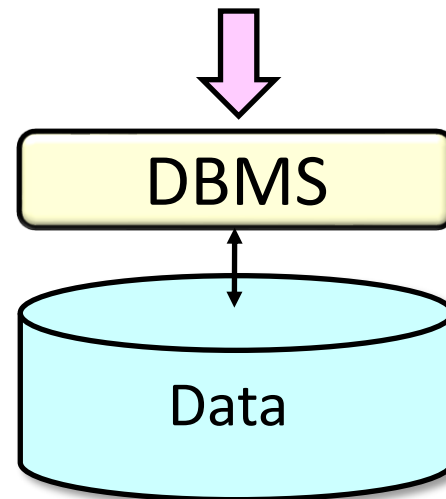


Resilience to System Failures

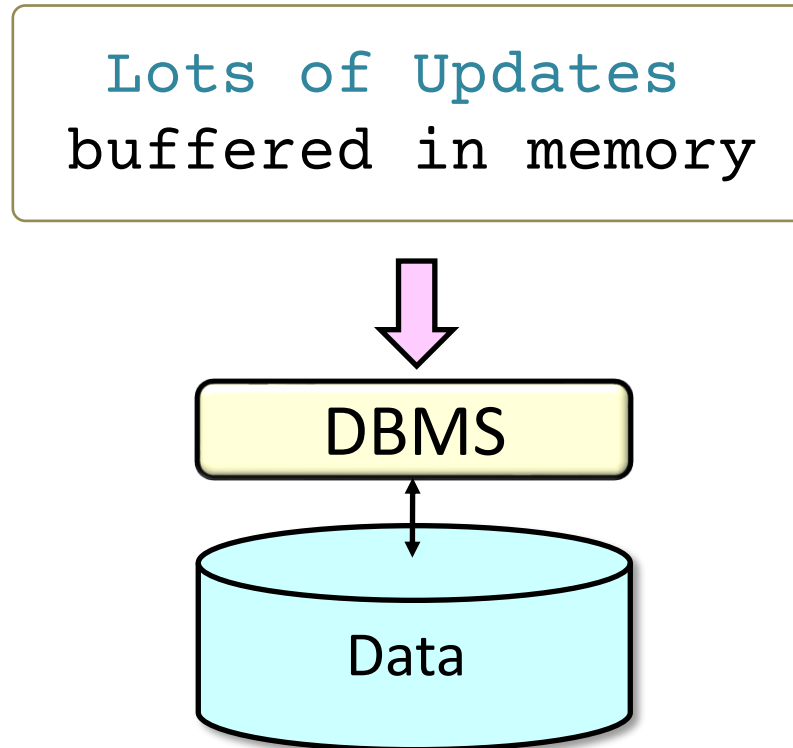


Resilience to System Failures

```
INSERT INTO Archive  
  SELECT * FROM Apply WHERE decision = 'N';  
DELETE FROM Apply WHERE decision = 'N';
```

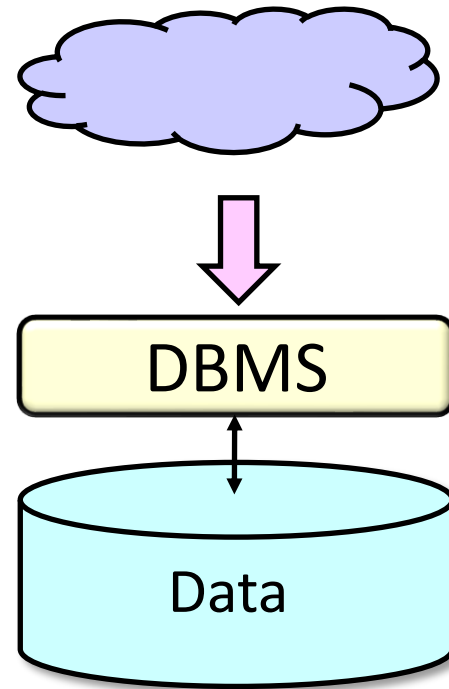


Resilience to System Failures



System-Failure Goal

Guarantee all-or-nothing execution, regardless of failures



Solution for both Concurrency and Failures

TRANSCATIONS

Guarantee all-or-nothing execution, regardless of failures

- Transactions appear to run in isolation
- If the system fails, each transaction's changes are reflected either entirely or not at all.



Solution for both Concurrency and Failures

TRANSCATIONS

A transaction is a sequence of one or more SQL operations treated as a unit.

SQL Standard:

- Transaction begins automatically on first SQL statement.
- On 'COMMIT' transaction ends and new one begins.
- Current transaction ends on session termination
- 'AUTOCOMMIT' turns each statement into transaction.



Solution for both Concurrency and Failures

TRANSCATIONS

A transaction is a sequence of one or more SQL operations treated as a unit.



ACID Properties

Atomicity

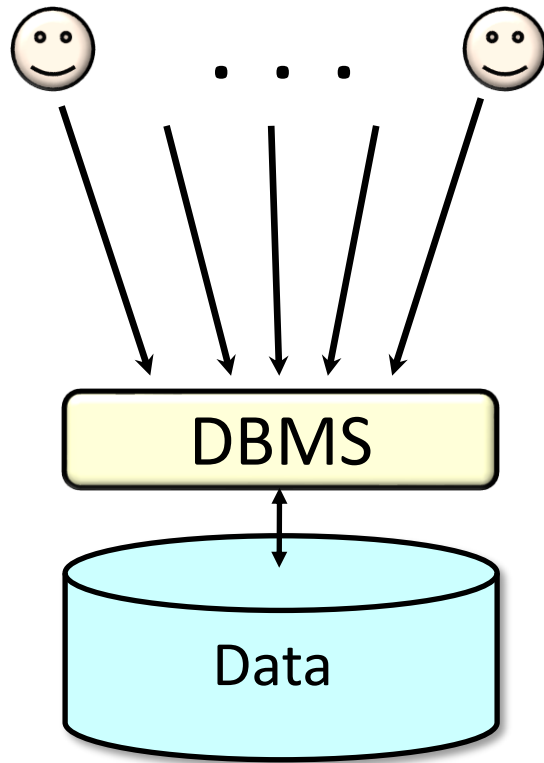
Consistency

Isolation

Durability



ACID Properties: Isolation



Serialisability

Operations may be interleaved, but execution must be equivalent to some sequential (serial) order of all transactions.

Locking



Concurrent Access: Attribute-Level Inconsistency

T1 `UPDATE College SET enrollment = enrollment + 1000
WHERE cName = 'University of London'`

... concurrent with ...

T2 `UPDATE College SET enrollment = enrollment + 1500
WHERE cName = 'University of London'`

T1; T2
or ...order doesn't matter...
T2; T1



Concurrent Access: Tuple-Level Inconsistency

```
UPDATE Apply SET major = 'CS'  
WHERE sID = 123
```

... concurrent with ...

```
UPDATE Apply SET decision = 'Y'  
WHERE sID = 123
```

T1; T2
or *...order doesn't matter...*
T2; T1



Concurrent Access: Table-Level Inconsistency

T1

```
UPDATE Apply SET decision = 'Y'
WHERE sid IN ( SELECT sid
                FROM Student
                WHERE GPA > 3.9 )
```

... concurrent with ...

T2

```
UPDATE Student SET GPA = (1.1)*GPA
WHERE sizeHS > 2500
```

T1; T2

or ...order matters...

T2; T1

*ACID only guarantees
isolation, not correctness.*



Concurrent Access: Multi-Statement Inconsistency

```
INSERT INTO Archive  
  SELECT * FROM Apply WHERE decision = 'N';  
DELETE FROM Apply WHERE decision = 'N';
```

... concurrent with ...

```
SELECT COUNT(*) FROM Apply;  
SELECT COUNT(*) FROM Archive;
```

T1; T2

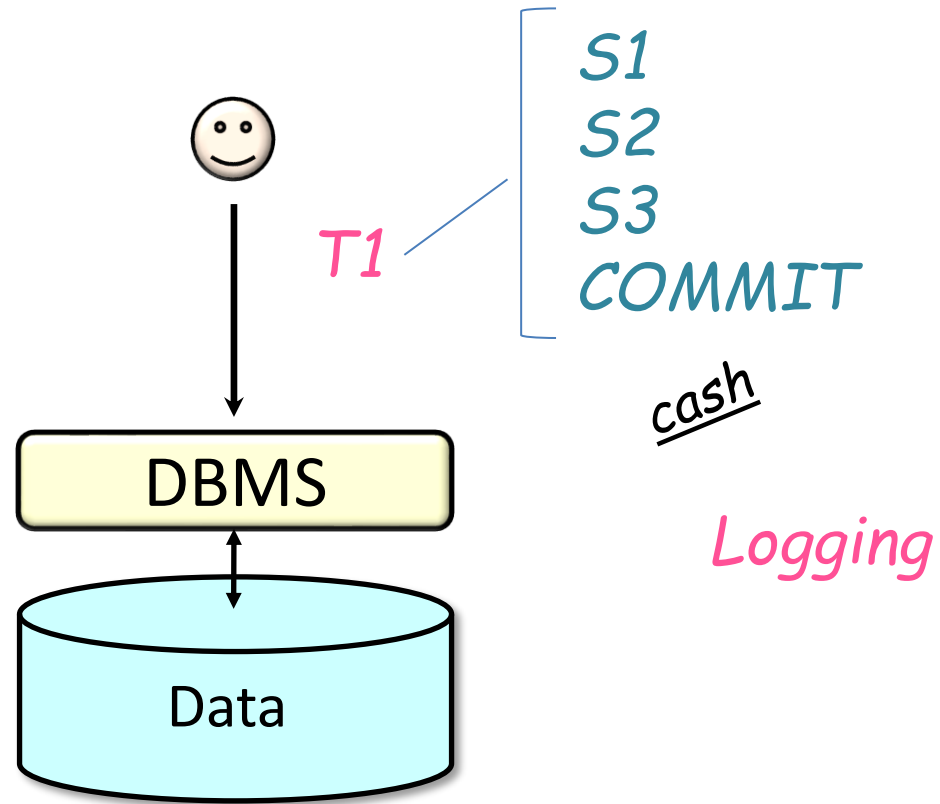
or ...order matters...

T2; T1

*ACID only guarantees
isolation, not correctness.*



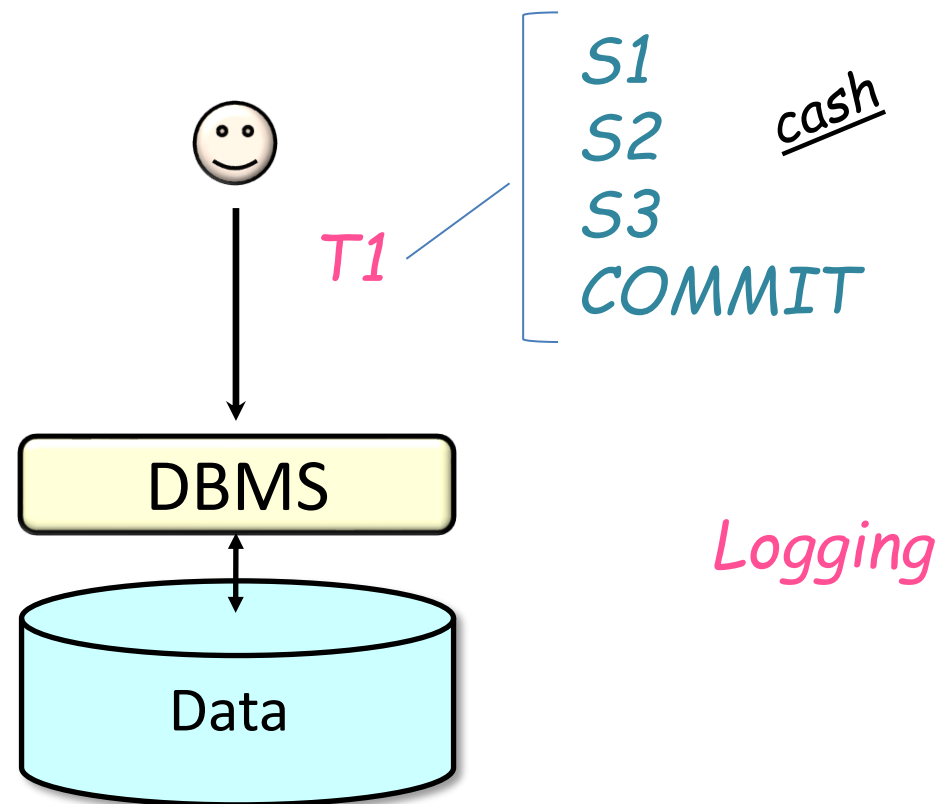
ACID Properties: Durability



If system crashes after transaction commits, all effects of the transaction remain in database.



ACID Properties: Atomicity



Each transaction is 'all-or-nothing' never left half done.



Transaction Rollback (=Abort)

- Undoes partial effects of transaction
- Can be system- or client-initiated

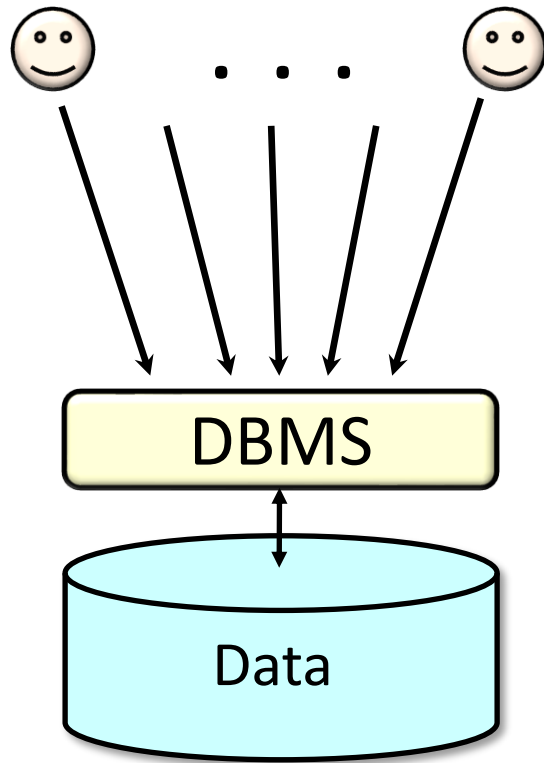
*Each transaction is 'all-or-nothing'
never left half done.*

```
Begin Transaction;  
<get input from user>  
SQL commands based on input  
<confirm results with user>  
If ans = 'ok Then Commit; Else Rollback;
```

Locking



ACID Properties: Consistency



Each Client, Each Transaction:

- We can assume all constraints hold when transaction begins.
- ACID must guarantee all constraints hold when transaction ends.

Serialisability → Constraints always hold.

Solution for both Concurrency and Failures

TRANSCATIONS

Atomicity

Consistency

Isolation

Durability



Solution for both Concurrency and Failures

TRANSCATIONS

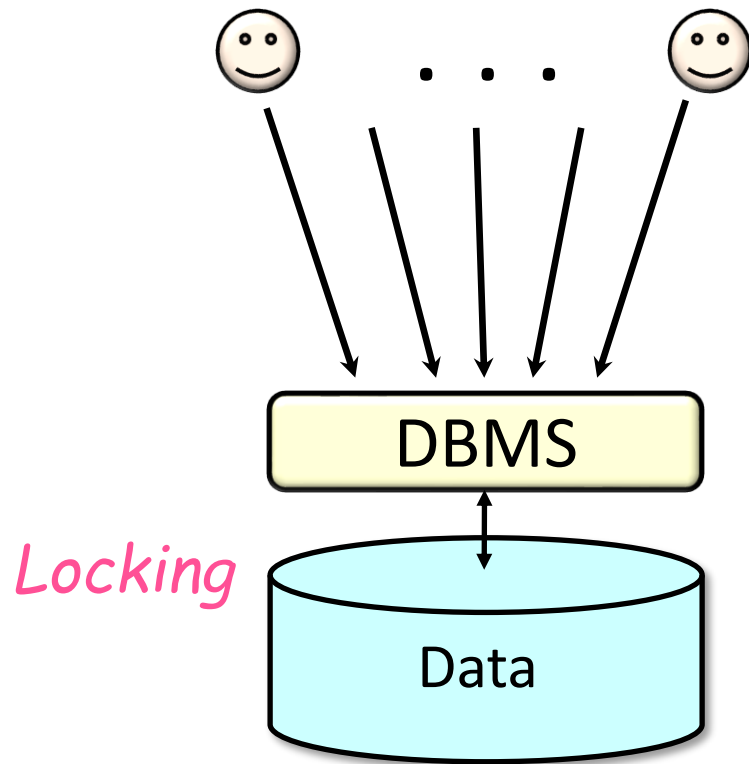
A transaction is a sequence of one or more SQL operations treated as a unit.

SQL Standard:

- Transactions appear to run in isolation
- If the system fails, each transaction's changes are reflected either entirely or not at all.



ACID Properties: Isolation



Serialisability

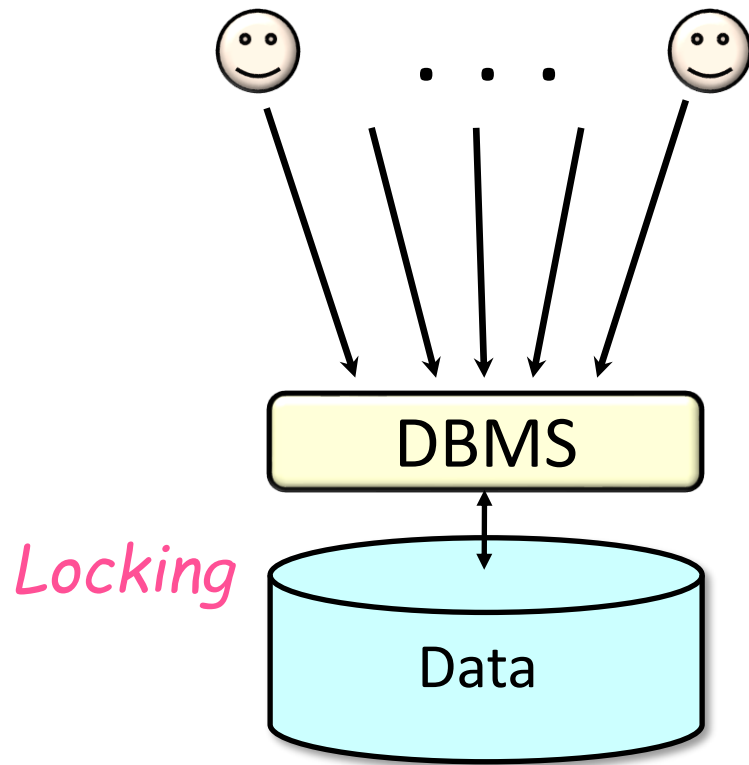
Operations may be interleaved, but execution must be equivalent to some sequential (serial) order of all transactions.

→ *Overhead*

→ *Reduction in Concurrency*



ACID Properties: Isolation



Weaker 'Isolation Levels'

Read Uncommitted
Read Committed
Repeatable Read

→ *Overhead*

→ *Reduction in Concurrency*

↓ *Overhead* ↑ *Concurrency*

↓ *Consistency Guarantees*



What is the practical usage of falsifying functional dependencies in sample tables?



Constraints and Triggers

- Relational Databases
- SQL Standard

(Integrity) Constraints

static

- constrain allowable database states

Triggers

- monitor database changes
- check conditions and initiate actions

dynamic



Integrity Constraints

Impose restrictions on allowable data, beyond those imposed by structure and types

Examples

`0.0 < GPA <= 4.0`

`decision : 'Y' 'N' NOT NULL`

`sizeHS < 200 → not admitted enrollment > 30,000`



Integrity Constraints

Impose restrictions on allowable data, beyond those imposed by structure and types

Why use them?

- Data-Entry Errors (for Inserts)
- Correctness Criteria (for Updates)
- Enforcing Consistency
- Tell system about data – store, query processing



Integrity Constraints

Impose restrictions on allowable data, beyond those imposed by structure and types

Classification

- NOT NULL
- Key (Primary Key)
- Referential Integrity (Foreign Key)ata-Entry Errors (for Inserts)
- Attribute-Based
- Tuple-Based
- General Assertions



Declaring and Enforcing Constraints

Declaration

- with original schema
- or later

Enforcement

- check after every modification
- deferred constraint checking



Triggers

Event-Condition-Action Rules

When event occurs, check condition; if true, do action

Examples

- enrollment > 35,000 → reject all applications
- insert application with GPA > 3.95 → accept automatically



Triggers

Event-Condition-Action Rules

When event occurs, check condition; if true, do action

Why use them?

- Move logic from applications into DBMS itself
- To enforce constraints
 - Expressiveness
 - Constraint 'Repair' Logic



Trigger in SQL

Event-Condition-Action Rules

```
CREATE TRIGGER name  
BEFORE|AFTER|INSTEAD OF events  
[referencing-variables]  
[FOR EACH ROW]  
WHEN (condition)  
action
```



Triggers

Event-Condition-Action Rules

When event occurs, check condition; if true, do action

- Move monitoring logic from apps into DBMS
- Enforce constraints
 - Beyond what constraint system supports
 - Automatic constraint 'repair'



Tricky Issues

- Row-Level vs Statement-Level
 - New/Old Row and New/Old Table
 - Before, Instead of
- Multiple triggers activated at the same time
- Trigger actions activating other triggers (chaining)
 - Also self-triggering, cycles, nested invocations
 - monitoring logic from apps into DBMS
- Conditions in **WHEN** vs as part of **ACTION**



Constraints and Triggers

- Relational Databases
- SQL Standard *systems vary considerably*

(Integrity) Constraints

- constrain allowable database states

Triggers

- monitor database changes
- check conditions and initiate actions

