

# Databases and Advanced Data Techniques

Tarapong Sreenuch

Basics of NoSQL

# Overview of NoSQL



## What is NoSQL?

- The NoSQL name was introduced during an open-source event on distributed databases
- It doesn't mean 'No SQL' - it means 'Not only SQL'

NoSQL



**'Not only  
SQL'**



## What is NoSQL?

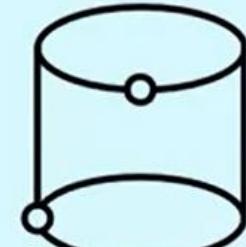
- Refers to family of databases that vary widely in style and technology
- However, they share a common trait
  - Non-relational
  - Not standard row and column type RDBMS
- Could be referred to as 'Non-relational databases'



## What is NoSQL?

### NoSQL databases:

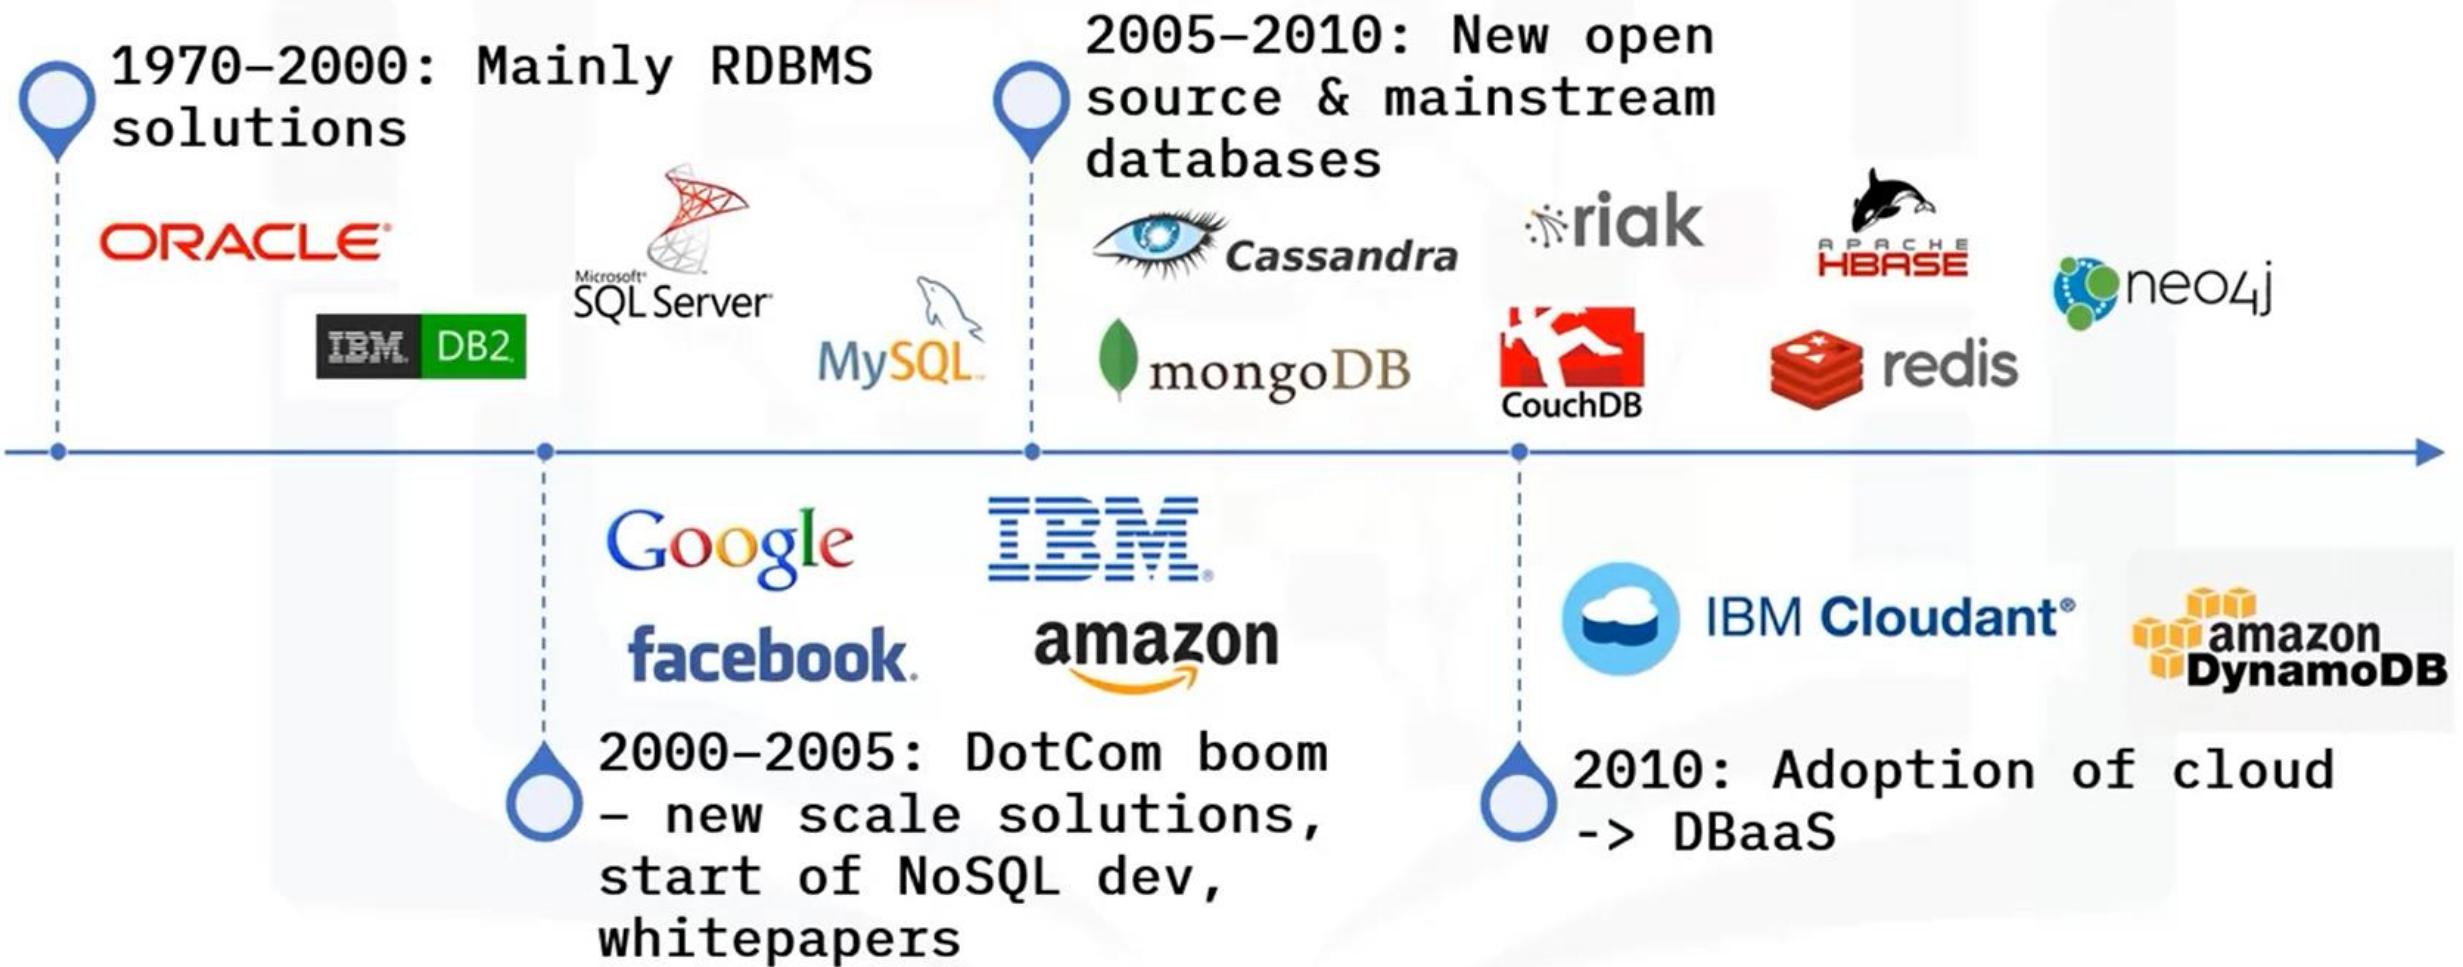
- Provide new ways of storing and querying data
  - Address several issues for modern apps
- Are designed to handle a different breed of scale - 'Big Data'
- Typically address more specialized use cases
- Simpler to develop app functionality for than RDBMS



NoSQL  
Databases



# History of NoSQL



## Recap: Overview of NoSQL

- The name 'NoSQL' stands for Not only SQL
- NoSQL refers to a class of databases that are non-relational in architecture
- Implementations of NoSQL databases differ technically, but share common traits
- Since 2000 NoSQL databases have become more popular in the marketplace, due to scale demands of Big Data



# Characteristics of NoSQL Databases



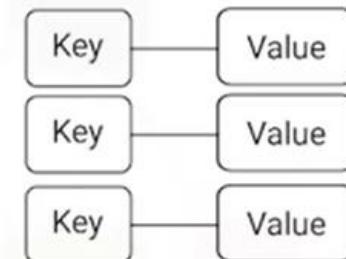
## NoSQL Database Categories

- The most common trait amongst NoSQL databases is that they are non-relational in architecture
- What types of NoSQL databases are available?
- What is common to them?



## NoSQL Database Categories

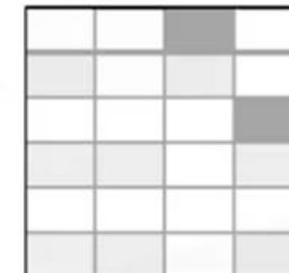
General consensus is...  
...NoSQL databases fit  
into four categories



Key-Value



Document



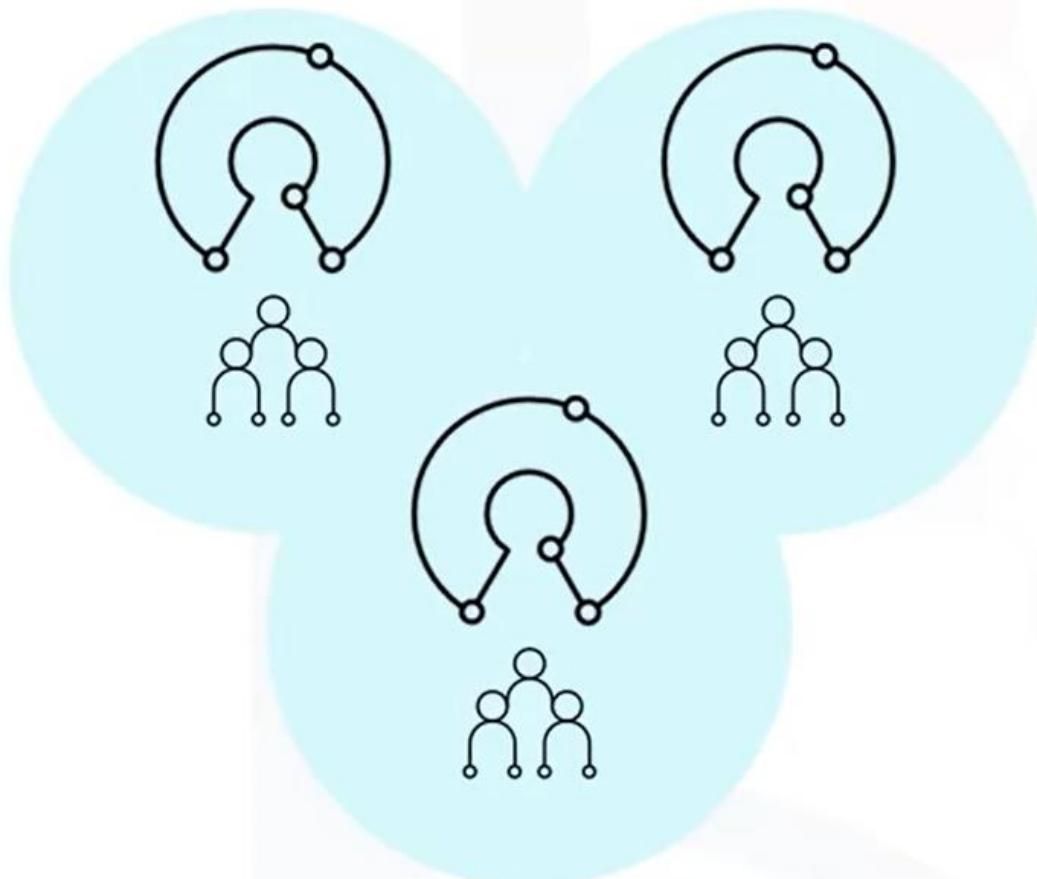
Column



Graph



## NoSQL Database Categories



Open Source  
Communities

But what ties NoSQL databases together?

- Majority have their roots in the open source community
- Many have been used and leveraged in an open source manner
- Open source community support is fundamental to their industry growth



## NoSQL Database Categories

- Companies often develop a commercial version of the database alongside the open source version
- Examples include:



**IBM Cloudant®**



**DATASTAX**



 **mongoDB**



## NoSQL Database Categories

- They all differ technically speaking, but have a few commonalities
- Most NoSQL databases:
  - Are built to scale horizontally
  - Share data more easily than RDBMS
  - Use a global unique key to simplify data sharding
  - Are more use case specific than RDBMS
  - Are more developer friendly than RDBMS
  - Allow more agile development via flexible schemas



## Benefits of NoSQL Databases

Why use a NoSQL database?

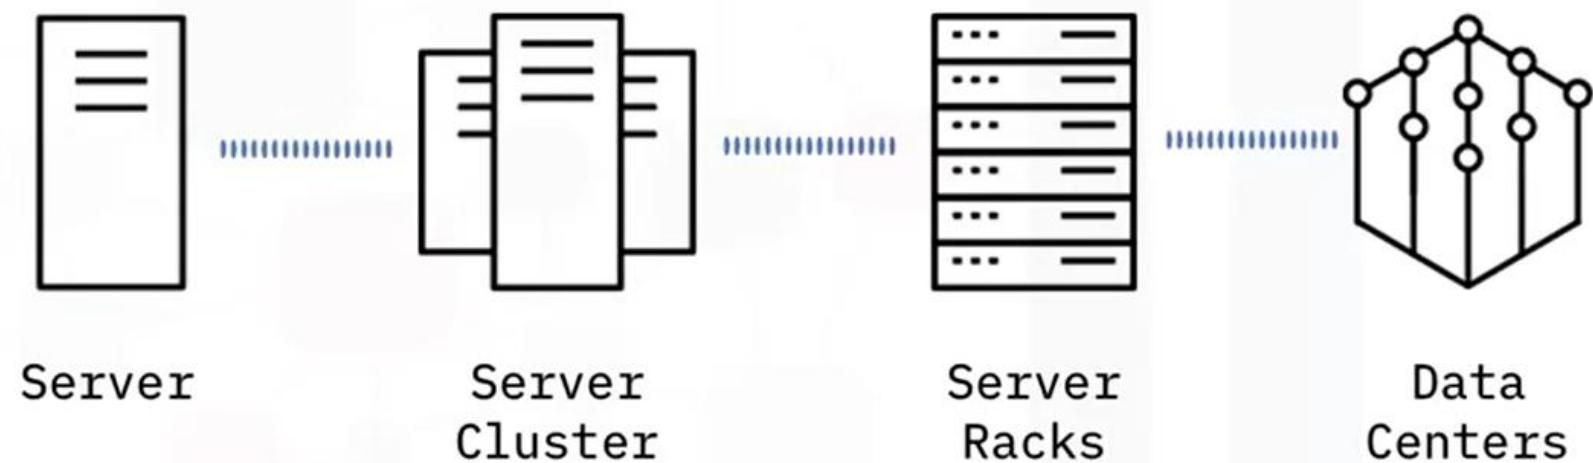


Why is their popularity growing so rapidly?



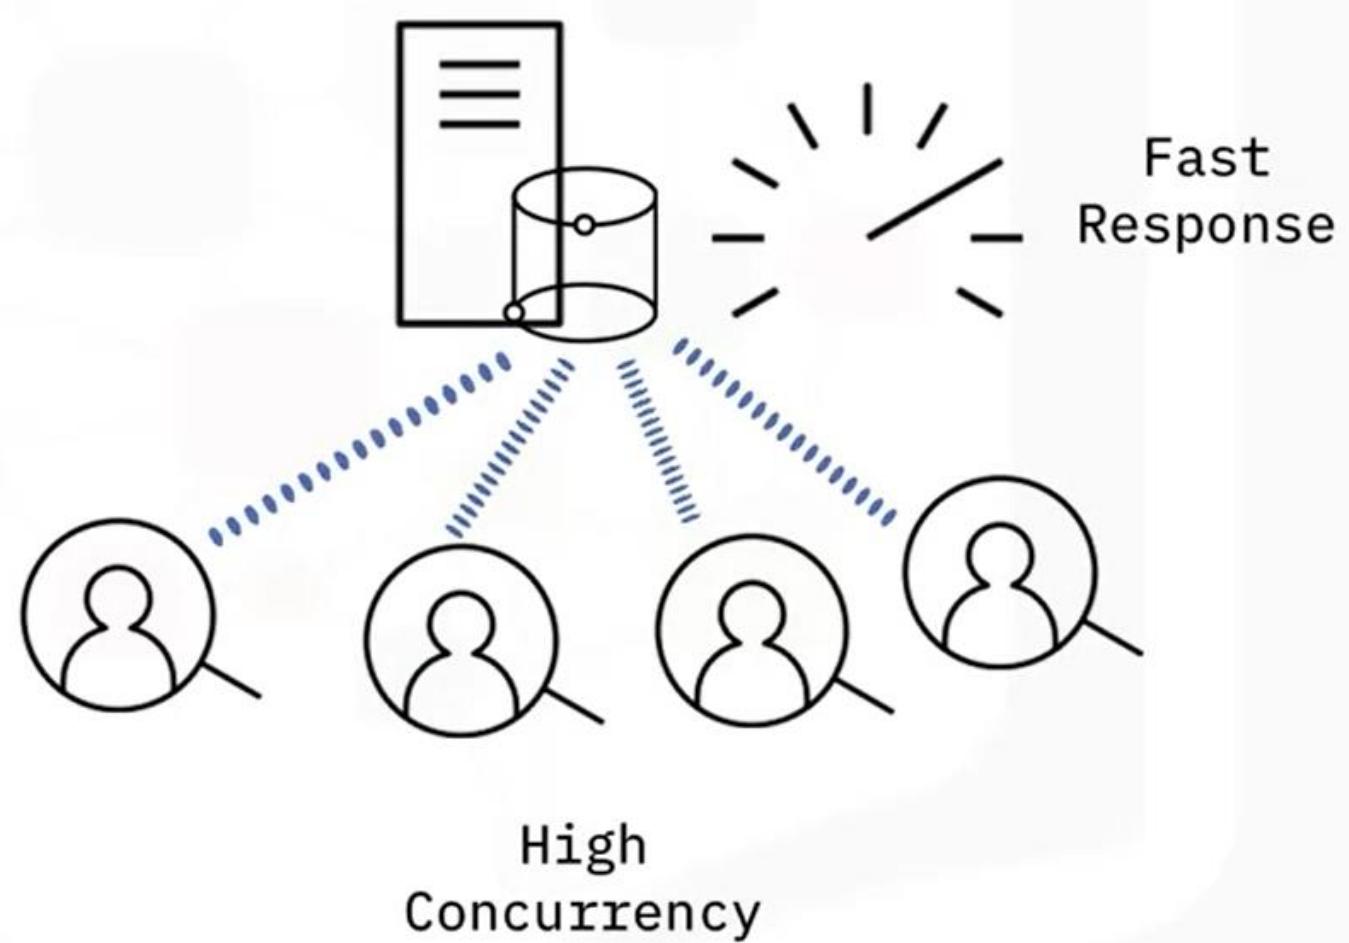
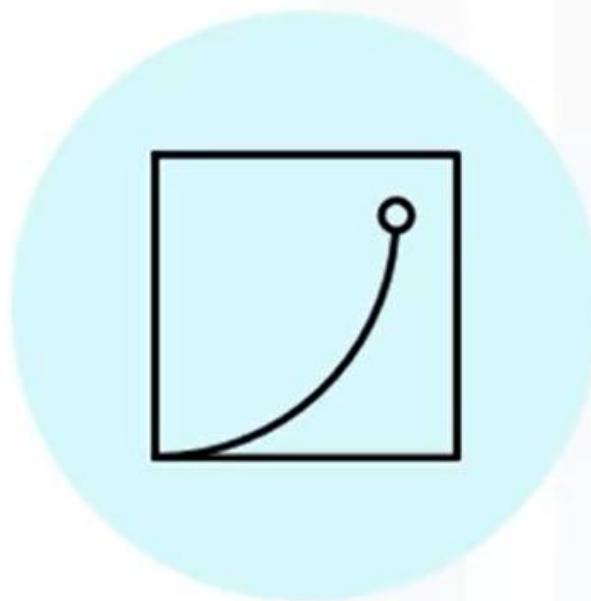
# Benefits of NoSQL Databases

## Scalability



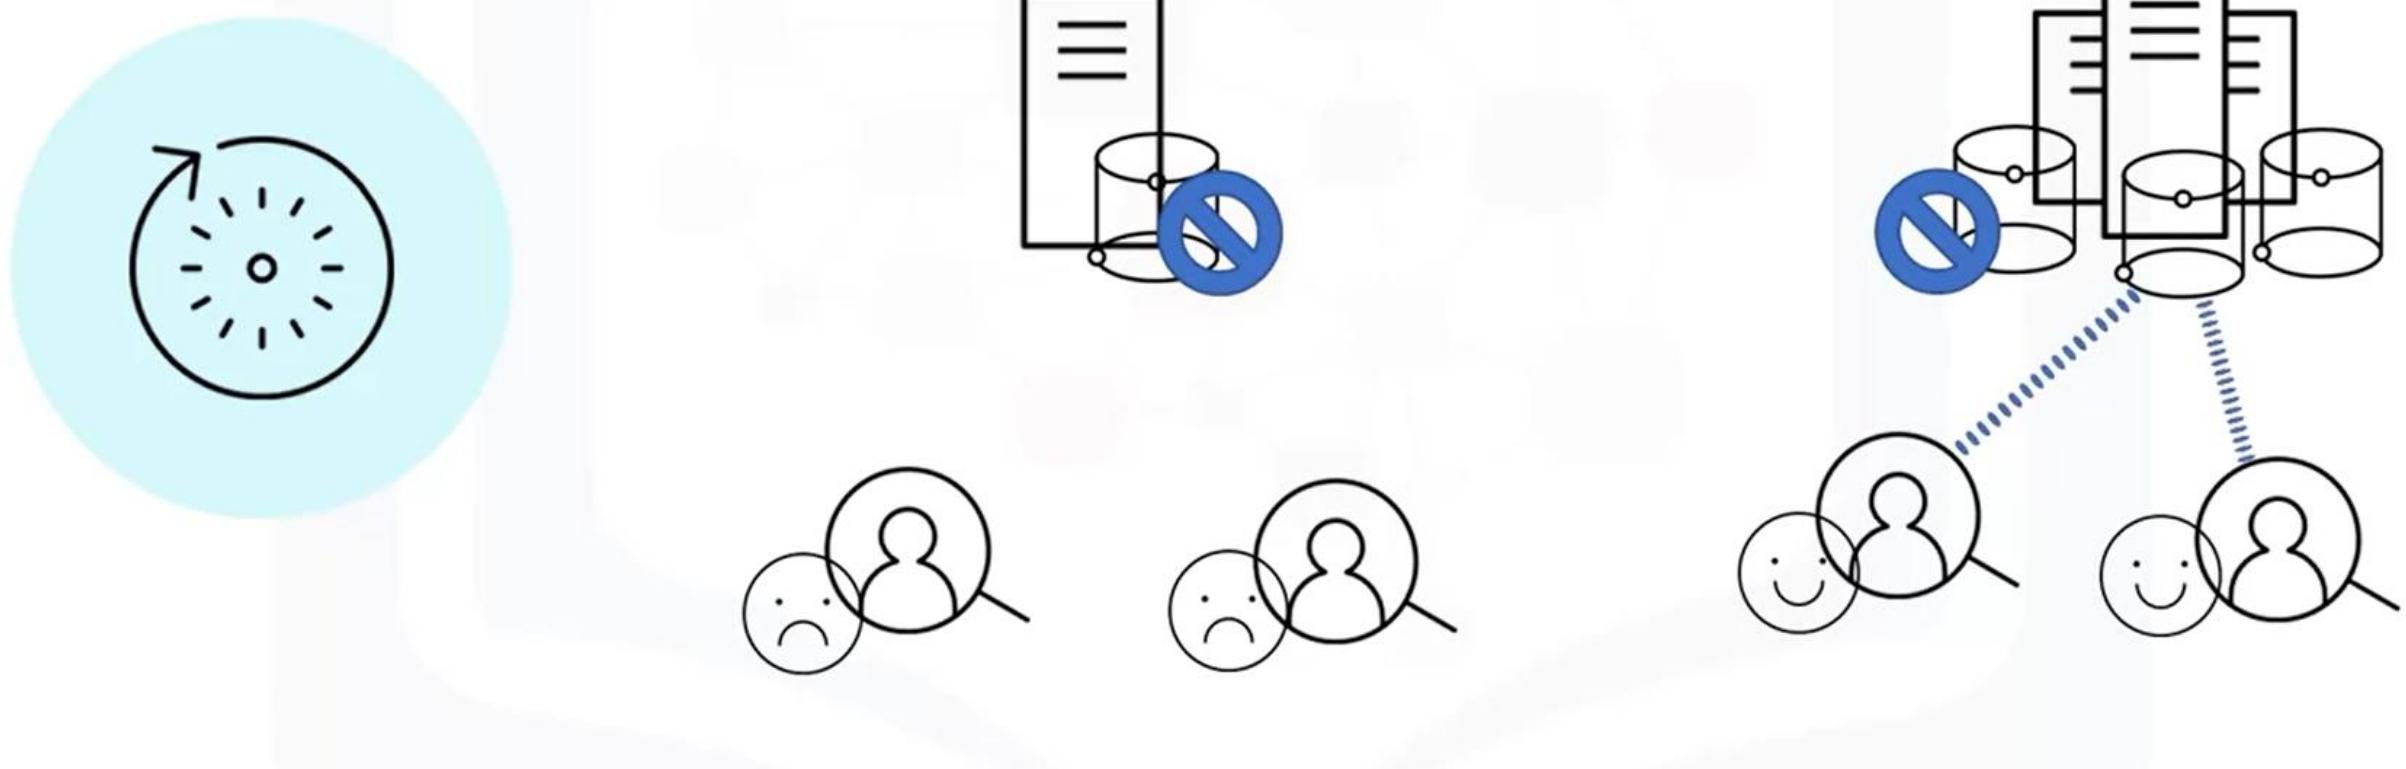
## Benefits of NoSQL Databases

### Performance



# Benefits of NoSQL Databases

## Availability



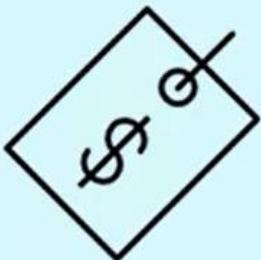
## Benefits of NoSQL Databases

### Cloud Architecture

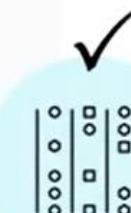
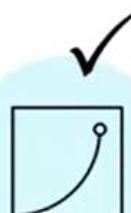
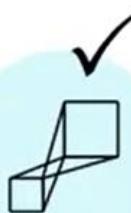


# Benefits of NoSQL Databases

Cost

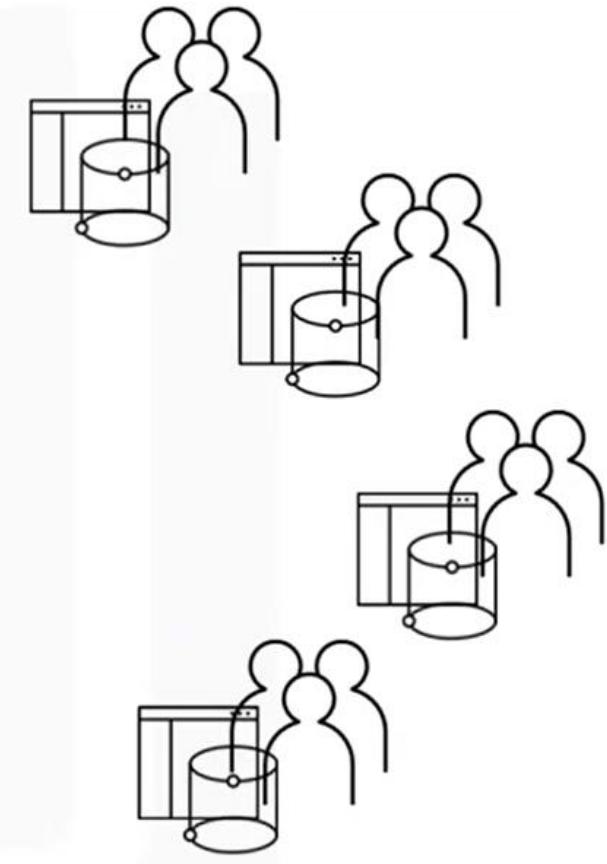
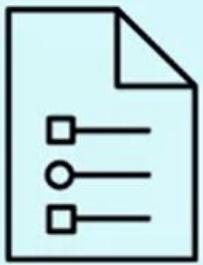


NoSQL



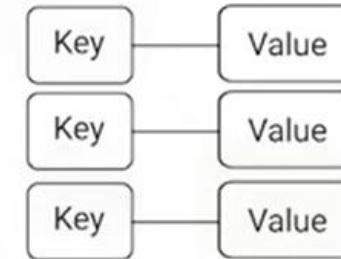
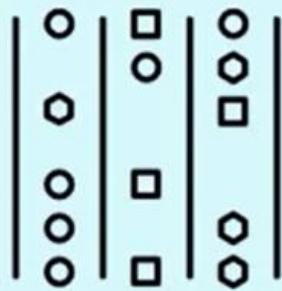
# Benefits of NoSQL Databases

## Flexible Schema

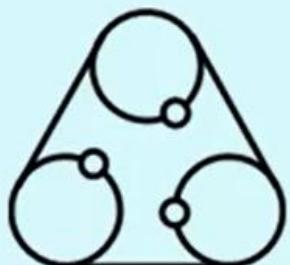


## Benefits of NoSQL Databases

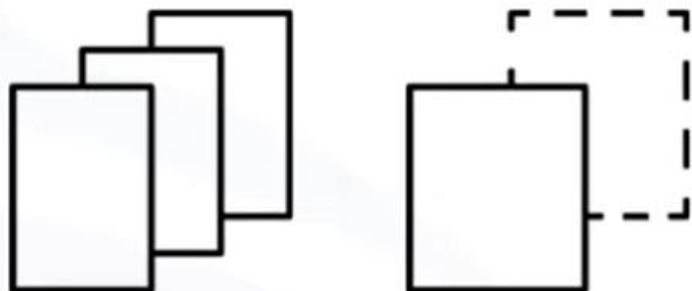
### Varied Data Structures



## Specialized Capabilities



Indexing and Querying



Data Replication Robustness



---



---

Modern HTTP APIs

## Benefits of NoSQL Databases

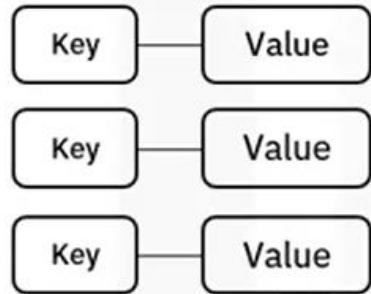
- NoSQL databases are non-relational
- There are four categories of NoSQL database
- NoSQL databases have their roots in the open-source community
- NoSQL database implementations are technically different
- There are several benefits to adopting NoSQL databases



# Key-Value NoSQL Databases



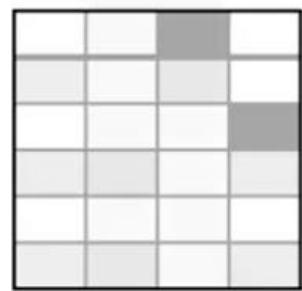
# NoSQL Database Categories



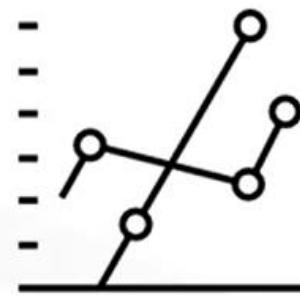
Key-Value



Document



Column

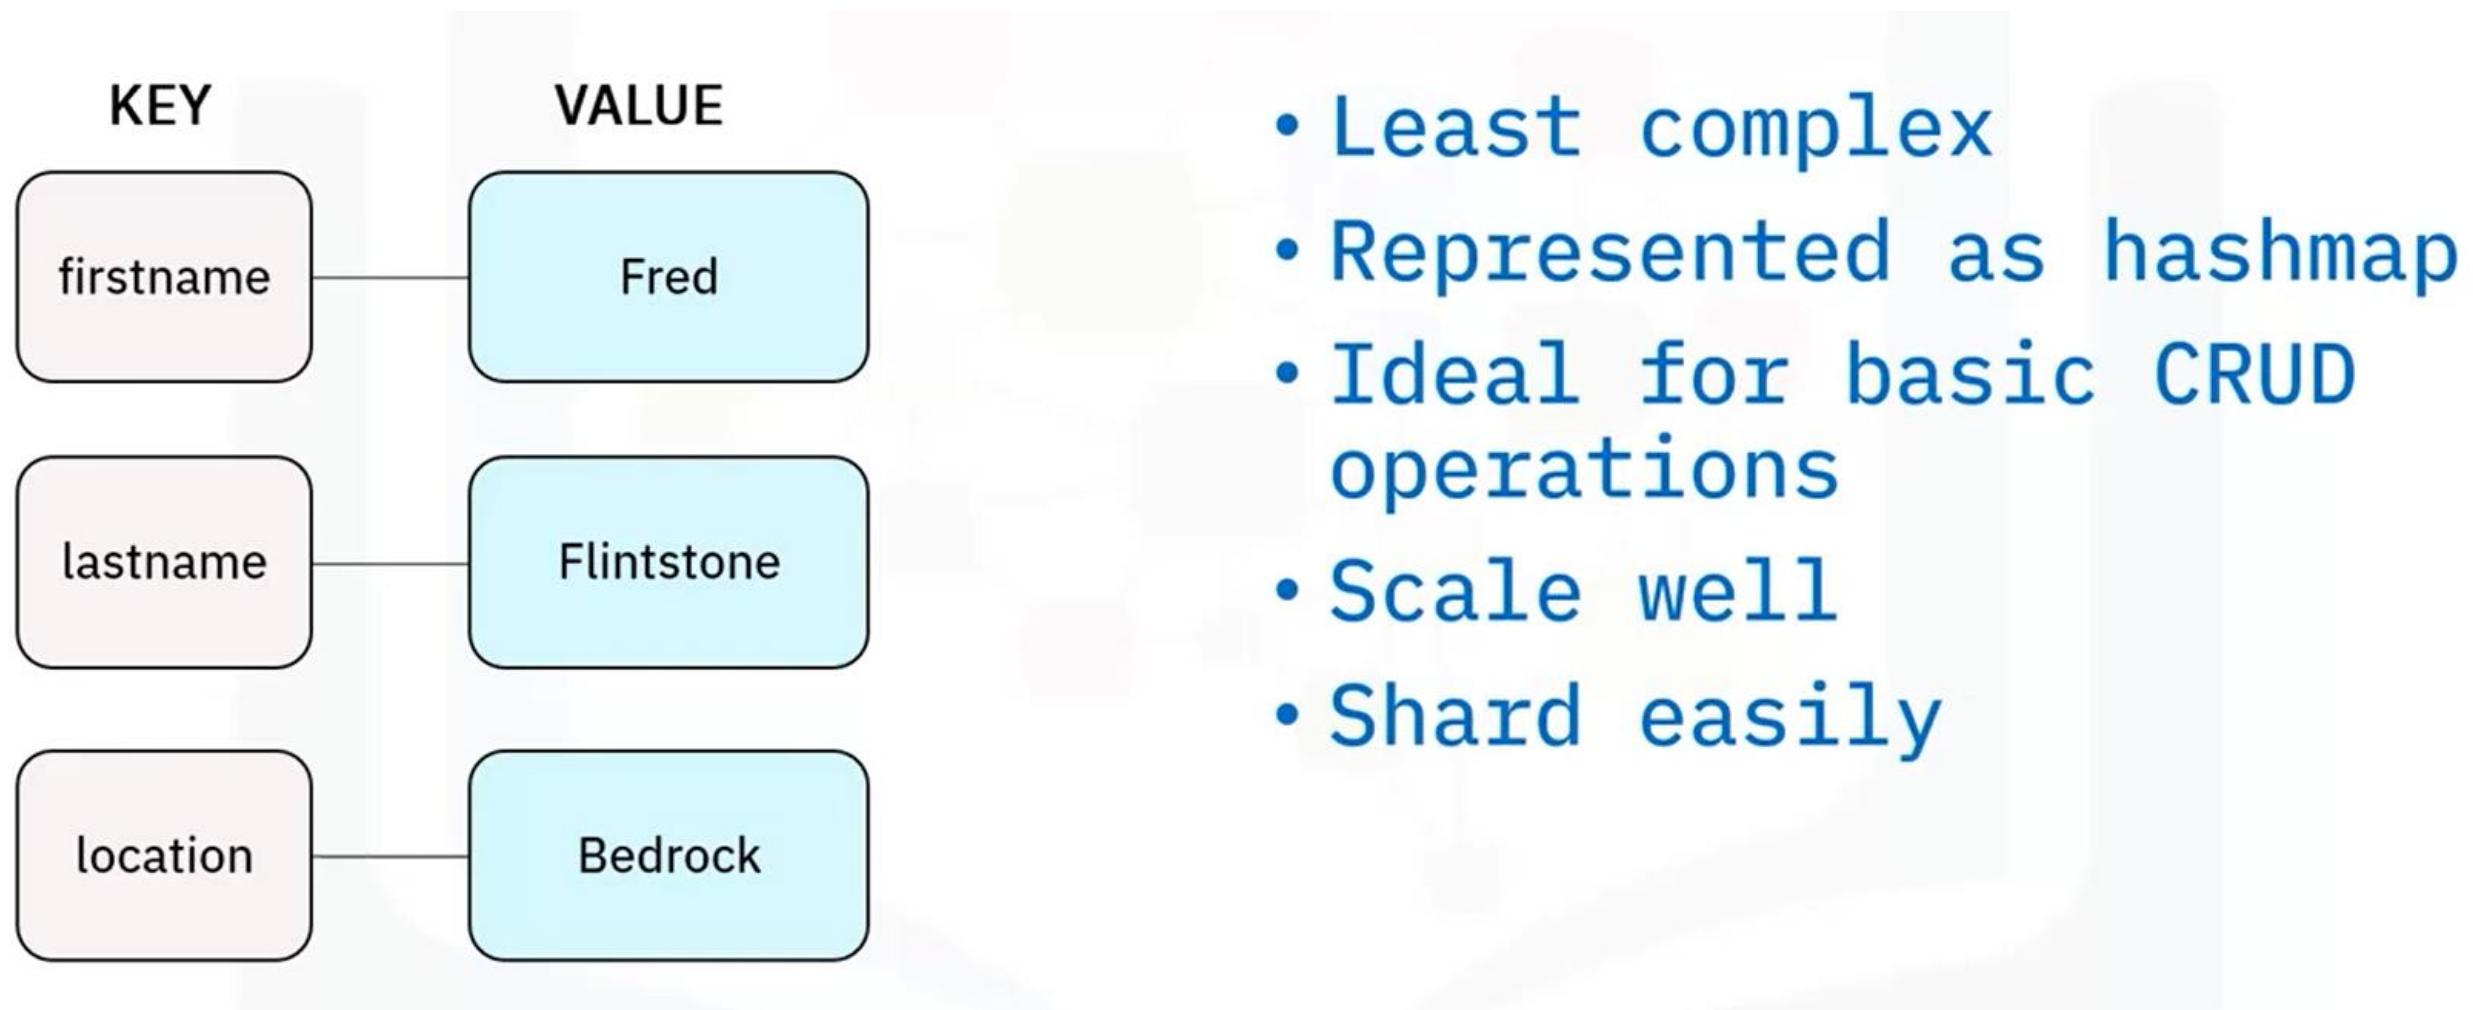


Graph

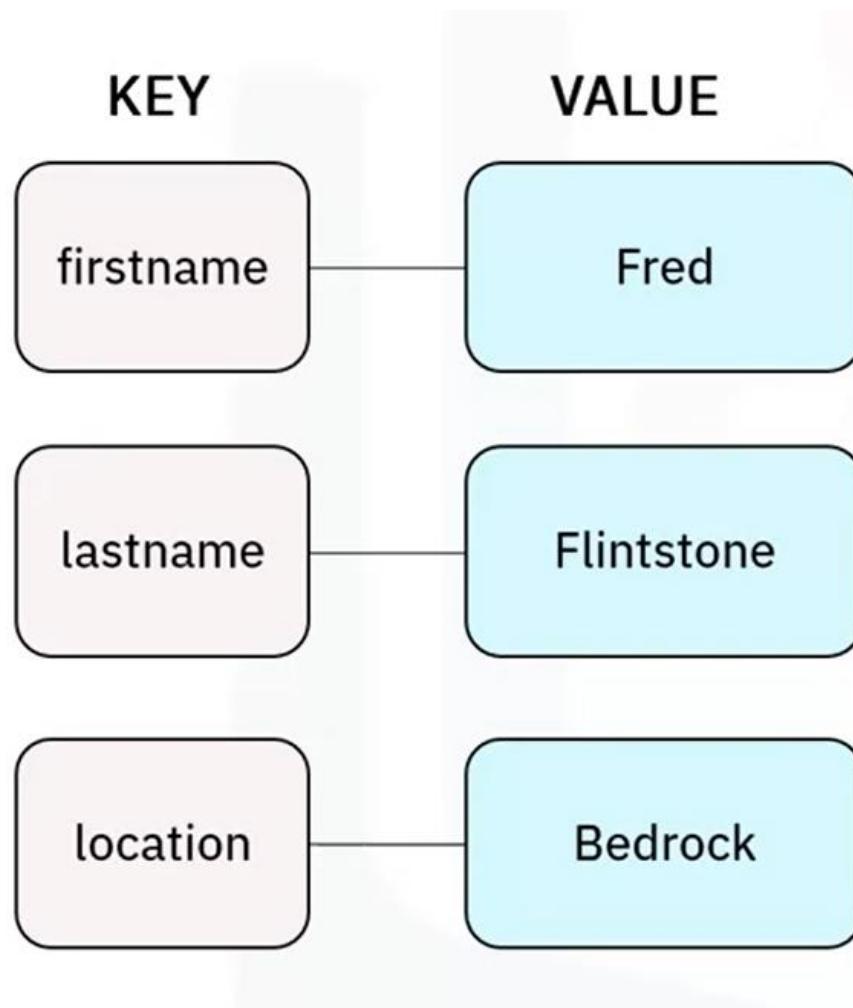
Each category has:

- Unique characteristics
- Architecture
- Use cases

# Key-Value NoSQL Database Architecture



# Key-Value NoSQL Database Architecture



- Not intended for complex queries
- Atomic for single key operations only
- Value blobs are opaque to database
  - Less flexible data indexing and querying



## Suitable Use Cases

- For quick basic CRUD operations on non-interconnected data
  - E.g. Storing and retrieving session information for web applications
- Storing in-app user profiles and preferences
- Shopping cart data for online stores

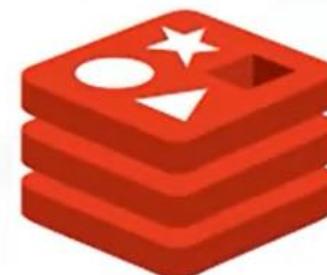


### Unsuitable Use Cases

- For data that is interconnected with many-to-many relationships
  - Social networks
  - Recommendation engines
- When high-level consistency is required for multi-operation transactions with multiple keys
  - Need a database that provides ACID transactions
- When apps runs queries based on value vs key
  - Consider using 'Document' category of NoSQL database

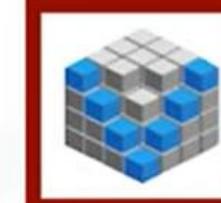


## Key-Value NoSQL Database Examples



redis

AEROSPIKE



Project Voldemort  
*A distributed database.*

## Recap: Key-Value NoSQL Databases

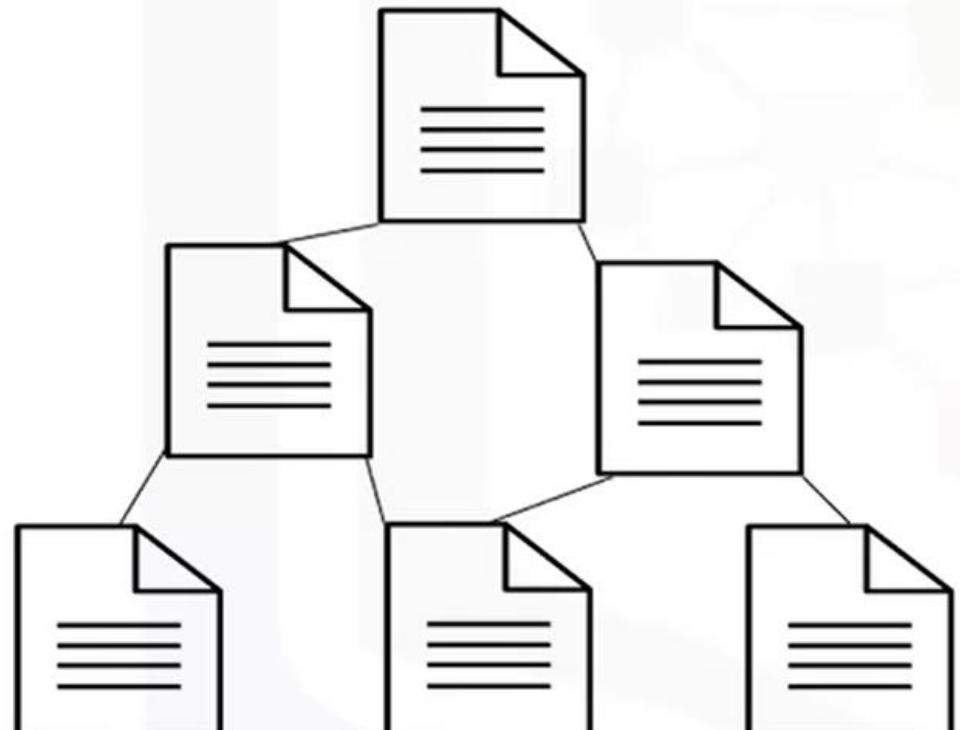
- The four main categories of NoSQL database are Key-Value, Document, Column, and Graph
- The Key-Value database architecture is the least complex; data is stored with a key and corresponding value blob and is represented by a hashmap
- The primary use cases for Key-Value databases are for quick CRUD operations; for example, storing and retrieving session information, storing in-app user profiles, and storing shopping cart data



# Document-Based NoSQL Databases



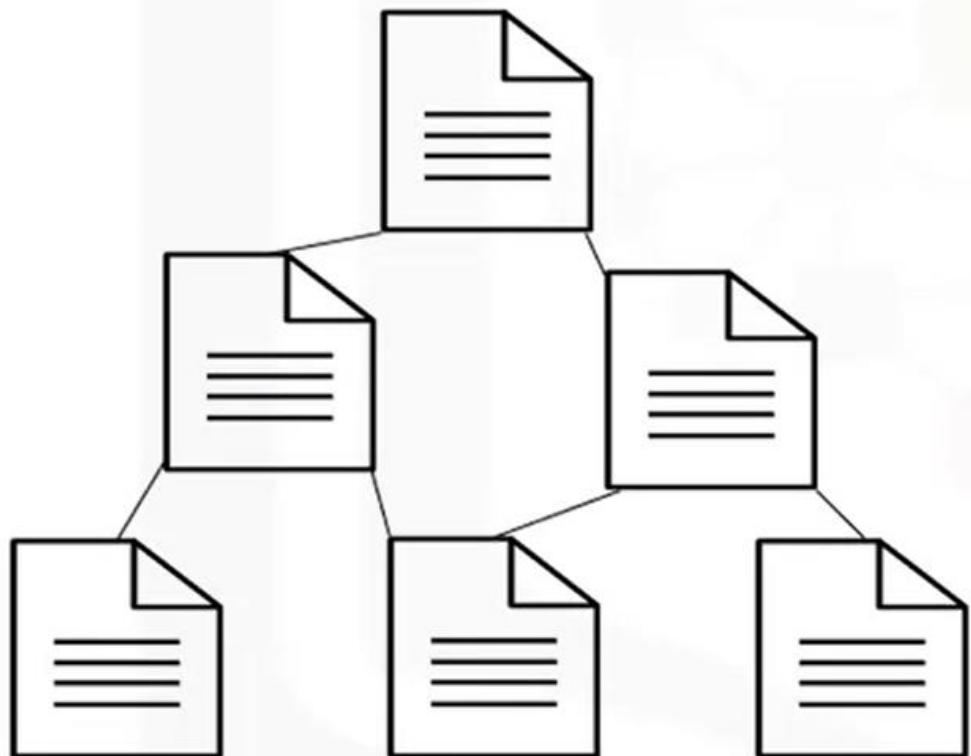
## Document-Based NoSQL Database Architecture



- Values are visible and can be queried
- Each piece of data is considered a document
  - Typically JSON or XML
- Each document offers a flexible schema
  - No two documents need to contain the same information



## Document-Based NoSQL Database Architecture



- Content of document databases can be indexed and queried
  - Key and value range lookups and search
  - Analytical queries with MapReduce
- Horizontally scalable
- Allow sharding across multiple nodes
- Typically only guarantee atomic operations on single documents



### Suitable Use Cases

- Event logging for apps and processes - each event instance is represented by a new document
- Online blogs - each user, post, comment, like, or action is represented by a document
- Operational datasets and metadata for web and mobile apps - designed with Internet in mind (JSON, RESTful APIs, unstructured data)



### Unsuitable Use Cases

- When you require ACID transactions
  - Document databases can't handle transactions that operate over multiple documents
  - Relational database would be a better choice
- If your data is in an aggregate-oriented design
  - If data naturally falls into a normalized tabular model
  - Relational database would be a better choice





IBM Cloudant®



CouchDB



mongoDB



RAVENDB

Couchbase



## Recap: Document-Based NoSQL Databases

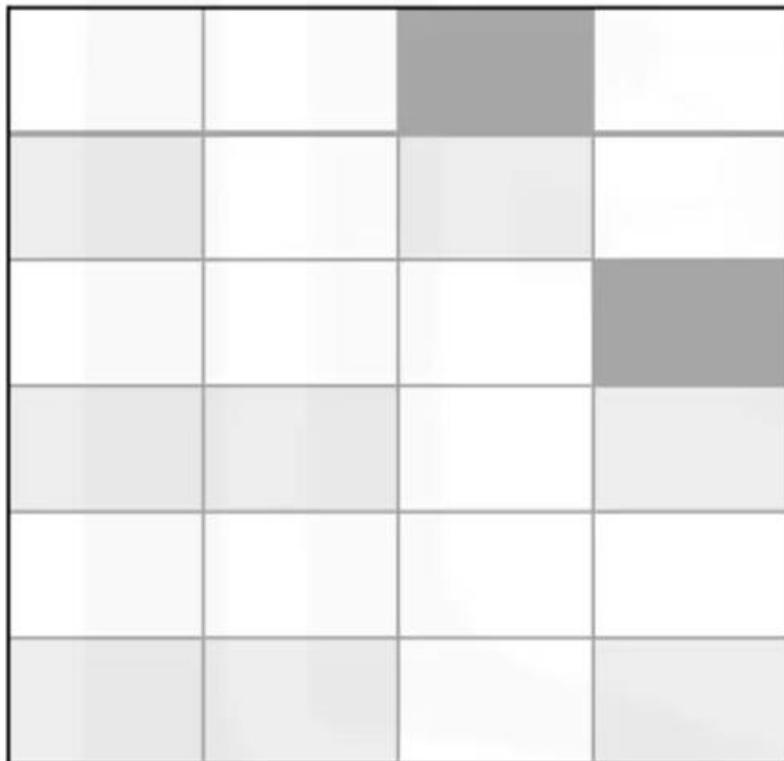
- Document-based NoSQL databases use documents to make values visible and able to be queried
- Each piece of data is considered a document (JSON/XML)
- Each document offers a flexible schema
- The primary use cases for document-based NoSQL databases are event logging for apps/processes, online blogging, operational datasets or metadata for web and mobile apps



# Column-Based NoSQL Databases



# Column-Based NoSQL Database Architecture



- Spawned from Google's 'Bigtable'
- a.k.a. Bigtable clones or Columnar or Wide-Column databases
- Store data in columns or groups of columns



## Column-Based NoSQL Database Architecture


- Column 'families' are several rows, with unique keys, belonging to one or more columns
  - Grouped in families as often accessed together
- Rows in a column family are not required to share the same columns
  - Can share all, a subset, or none
  - Columns can be added to any number of rows, or not



## Column-Based NoSQL Database Use Cases

### Suitable Use Cases

- Great for large amounts of sparse data
- Column databases can handle being deployed across clusters of nodes
- Column databases can be used for event logging and blogs
- Counters are a unique use case for column databases
- Columns can have a TTL parameter, making them useful for data with an expiration value



### Unsuitable Use Cases

- For traditional ACID transactions
  - Reads and writes are only atomic at the row level
- In early development, query patterns may change and require numerous changes to column-based databases
  - Can be costly and can slow down the production timeline



## Column-Based NoSQL Database Examples



## Recap: Column-Based NoSQL Databases

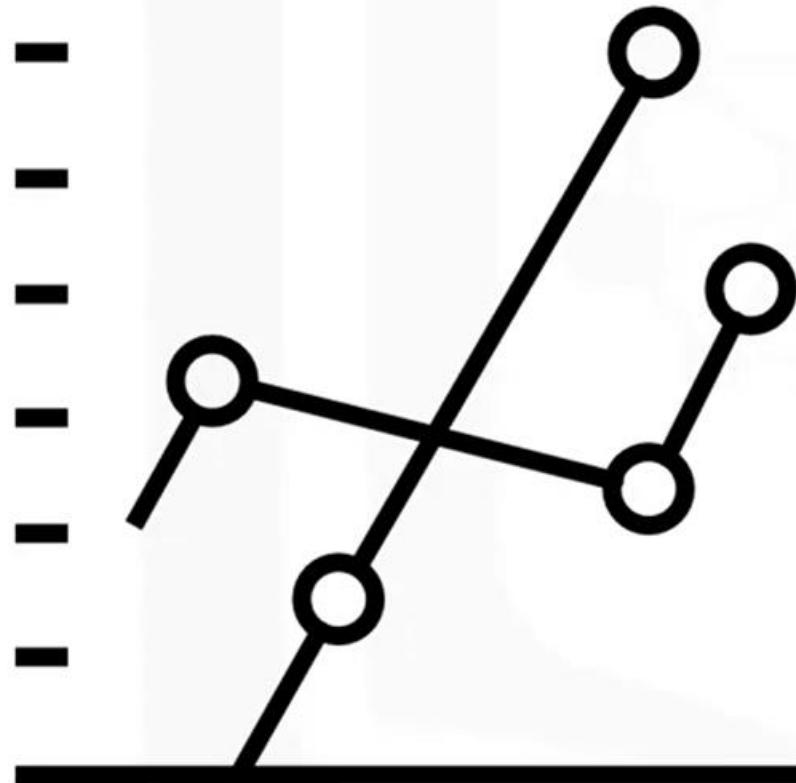
- Column-based databases were spawned from the architecture of Google's Bigtable storage system
- Column-based databases store data in columns or groups of columns
- Column 'families' are several rows, with unique keys, belonging to one or more columns
- The primary use cases for Column-based databases are event logging, blogs, counters, and data with expiration values



# Graph NoSQL Databases



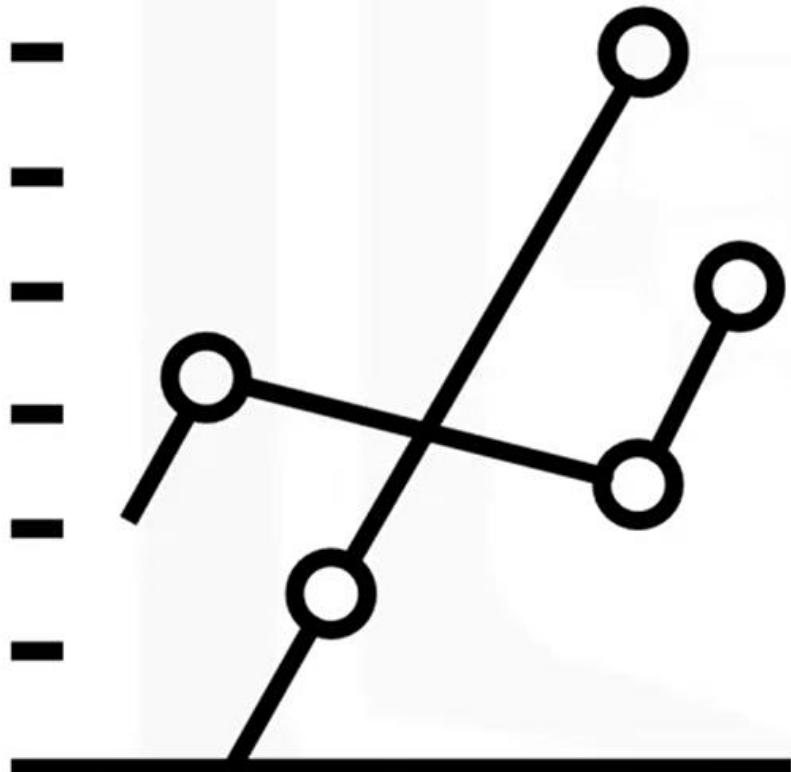
## Graph NoSQL Database Architecture



- Graph databases store information in entities (or nodes), and relationships (or edges)
- Graph databases are impressive when your data set resembles a graph-like data structure



## Graph NoSQL Database Architecture



- Graph databases do not shard well
  - Traversing a graph with nodes split across multiple servers can become difficult and hurt performance
- Graph databases are ACID transaction compliant
  - Unlike other NoSQL databases discussed



### Suitable Use Cases

- For highly connected and related data
- Social networking
- Routing, spatial, and map apps
- Recommendation engines



### Unsuitable Use Cases

- When looking for advantages offered by other NoSQL database categories
- When an application needs to scale horizontally
  - You will quickly reach the limitations associated with these types of data stores
- When trying to update all or a subset of nodes with a given parameter
  - These types of operations can prove to be difficult and non-trivial





## Recap: Graph NoSQL Databases

- Graph databases store information in entities and relationships
- Graph databases are impressive when your data set resembles a graph-like data structure
- Graph databases don't shard well but are ACID transaction compliant
- The primary use cases for Graph databases are for highly connected and related data, for social networking sites, for routing, spatial and map applications, and for recommendation engines

