"Active" data

- Constraint enforcement: When an operation violates a constraint, abort the operation or try to "fix" data
  - Example: enforcing referential integrity constraints
  - Generalize to arbitrary constraints?

- Data monitoring: When something happens to the data, automatically execute some action
  - Example: When price rises above $20 per share, sell
  - Example: When enrollment is at the limit and more students try to register, email the instructor

Triggers

- A trigger is an event-condition-action (ECA) rule
  - When event occurs, test condition; if condition is satisfied, execute action

- Example:
  - Event: some user's popularity is updated
  - Condition: the user is a member of "Jessica's Circle," and *pop* drops below 0.5
  - Action: kick that user out of Jessica's Circle

http://pt.simpsons.wikia.com/wiki/Arquivo:Jessica_lovejoy.jpg

Trigger example

```
CREATE TRIGGER PickyJessica
AFTER UPDATE OF pop ON User          Event
REFERENCING NEW ROW AS newUser
FOR EACH ROW                         Condition
WHEN (newUser.pop < 0.5)
 AND (newUser.uid IN (SELECT uid
                      FROM Member
                      WHERE gid = 'jes'))
DELETE FROM Member
WHERE uid = newUser.uid AND gid = 'jes';
                                     Action
```

Trigger options

- Possible events include:
  - INSERT ON *table*
  - DELETE ON *table*
  - UPDATE [OF *column*] ON *table*

- Granularity—trigger can be activated:
  - FOR EACH ROW modified
  - FOR EACH STATEMENT that performs modification

- Timing—action can be executed:
  - AFTER or BEFORE the triggering event
  - INSTEAD OF the triggering event on views (more later)

Transition variables

- OLD ROW: the modified row before the triggering event

- NEW ROW: the modified row after the triggering event

- OLD TABLE: a hypothetical read-only table containing all rows to be modified before the triggering event

- NEW TABLE: a hypothetical table containing all modified rows after the triggering event

☞Not all of them make sense all the time, e.g.

  - AFTER INSERT statement-level triggers
    - Can use only NEW TABLE
  - BEFORE DELETE row-level triggers
    - Can use only OLD ROW
  - etc.

Statement-level trigger example

```
CREATE TRIGGER PickyJessica
AFTER UPDATE OF pop ON User
REFERENCING NEW TABLE AS newUsers
FOR EACH STATEMENT
DELETE FROM Member
WHERE gid = 'jes'
AND uid IN (SELECT uid
            FROM newUsers
            WHERE pop < 0.5);
```

BEFORE trigger example

- Never allow age to decrease

```
CREATE TRIGGER NoFountainOfYouth
BEFORE UPDATE OF age ON User
REFERENCING OLD ROW AS o,
                NEW ROW AS n
FOR EACH ROW
WHEN (n.age < o.age)
SET n.age = o.age;
```

☞BEFORE triggers are often used to "condition" data

☞Another option is to raise an error in the trigger body to abort the transaction that caused the trigger to fire

Statement- vs. row-level triggers

Why are both needed?

- Certain triggers are only possible at statement level
  - If the number of users inserted by this statement exceeds 100 and their average age is below 13, then …

- Simple row-level triggers are easier to implement
  - Statement-level triggers require significant amount of state to be maintained in OLD TABLE and NEW TABLE
  - However, a row-level trigger gets fired for each row, so complex row-level triggers may be inefficient for statements that modify many rows

System issues

- Recursive firing of triggers
  - Action of one trigger causes another trigger to fire
  - Can get into an infinite loop
    - Some DBMS leave it to programmers/database administrators (e.g., PostgreSQL)
    - Some restrict trigger actions (e.g., Oracle)
    - Many set a maximum level of recursion (e.g., 16 in DB2)

- Interaction with constraints (very tricky to get right!)
  - When do we check if a triggering event violates constraints?
    - After a BEFORE trigger (so the trigger can fix a potential violation)
    - Before an AFTER trigger
  - AFTER triggers also see the effects of, say, cascaded deletes caused by referential integrity constraint violations
  - (Based on DB2; other DBMS may implement a different policy)

Views

- A view is like a "virtual" table
  - Defined by a query, which describes how to compute the view contents on the fly
  - DBMS stores the view definition query instead of view contents
  - Can be used in queries just like a regular table

Creating and dropping views

- Example: members of Jessica's Circle
  - `CREATE VIEW` JessicaCircle `AS`
    ```
    SELECT * FROM User
    WHERE uid IN (SELECT uid FROM Member
                         WHERE gid = 'jes');
    ```
  - Tables used in defining a view are called "base tables"
    - *User* and *Member* above

- To drop a view
  - `DROP VIEW` JessicaCircle;

Using views in queries

- Example: find the average popularity of members in Jessica's Circle
    - `SELECT AVG(pop) FROM JessicaCircle;`
    - To process the query, replace the reference to the view by its definition
    - ```
SELECT AVG(pop)
FROM (SELECT * FROM User
        WHERE uid IN
          (SELECT uid FROM Member
           WHERE gid = 'jes'))
       AS JessicaCircle;
```

# Why use views?

- To hide data from users

- To hide complexity from users

- Logical data independence
  - If applications deal with views, we can change the underlying schema without affecting applications
  - Recall physical data independence: change the physical organization of data without affecting applications

- To provide a uniform interface for different implementations or sources

☞Real database applications use tons of views

# Modifying views

- Does it even make sense, since views are virtual?

- It does make sense if we want users to really see views as tables

- Goal: modify the base tables such that the modification would appear to have been accomplished on the view

# A simple case

```
CREATE VIEW UserPop AS
SELECT uid, pop FROM User;
```

```
DELETE FROM UserPop WHERE uid = 123;
```

translates to:

```
DELETE FROM User WHERE uid = 123;
```

An impossible case

```
CREATE VIEW PopularUser AS
  SELECT uid, pop FROM User
  WHERE pop >= 0.8;
```

```
INSERT INTO PopularUser
  VALUES(987, 0.3);
```

- No matter what we do on *User*, the inserted row will not be in *PopularUser*

A case with too many possibilities

```
CREATE VIEW AveragePop(pop) AS
  SELECT AVG(pop) FROM User;
```
  - Note that you can rename columns in view definition

```
UPDATE AveragePop SET pop = 0.5;
```
- Set everybody's *pop* to 0.5?
- Adjust everybody's *pop* by the same amount?
- Just lower Jessica's *pop*?

SQL92 updateable views

- More or less just single-table selection queries
  - No join
  - No aggregation
  - No subqueries

- Arguably somewhat restrictive

- Still might get it wrong in some cases
  - See the slide titled "An impossible case"
  - Adding WITH CHECK OPTION to the end of the view definition will make DBMS reject such modifications

INSTEAD OF triggers for views

```
CREATE TRIGGER AdjustAveragePop
INSTEAD OF UPDATE ON AveragePop
REFERENCING OLD ROW AS o,
             NEW ROW AS n
FOR EACH ROW
UPDATE User
SET pop = pop + (n.pop-o.pop);
```

- What does this trigger do?

Indexes

- An index is an auxiliary persistent data structure
  - Search tree (e.g., B⁺-tree), lookup table (e.g., hash table), etc.

☞More on indexes later in this course!

- An index on $R.A$ can speed up accesses of the form
  - $R.A = value$
  - $R.A > value$ (sometimes; depending on the index type)
- An index on $(R.A_1, \dots, R.A_n)$ can speed up
  - $R.A_1 = value_1 \wedge \cdots \wedge R.A_n = value_n$
  - $(R.A_1, \dots, R.A_n) > (value_1, \dots, value_n)$ (again depends)

☞Ordering or index columns is important—is an index on $(R.A, R.B)$ equivalent to one on $(R.B, R.A)$?

☞How about an index on $R.A$ plus another on $R.B$?

Examples of using indexes

- `SELECT * FROM User WHERE name = 'Bart';`
  - Without an index on *User.name*: must scan the entire table if we store *User* as a flat file of unordered rows
  - With index: go "directly" to rows with `name='Bart'`
- `SELECT * FROM User, Member`
  `WHERE User.uid = Member.uid`
  `AND Member.gid = 'jes';`
  - With an index on *Member.gid* or (*gid, uid*): find relevant *Member* rows directly
  - With an index on *User.uid*: for each relevant *Member* row, directly look up *User* rows with matching *uid*
    - Without it: for each *Member* row, scan the entire *User* table for matching *uid*
      - Sorting could help

Creating and dropping indexes in SQL

CREATE [UNIQUE] INDEX $indexname$ ON $tablename$($columnname_1$,…,$columnname_n$);

- With UNIQUE, the DBMS will also enforce that $\{columnname_1, …, columnname_n\}$ is a key of $tablename$

DROP INDEX $indexname$;

- Typically, the DBMS will automatically create indexes for PRIMARY KEY and UNIQUE constraint declarations

Choosing indexes to create

More indexes = better performance?

- Indexes take space

- Indexes need to be maintained when data is updated

- Indexes have one more level of indirection

☞Optimal index selection depends on both query and update workload and the size of tables

- Automatic index selection is now featured in some commercial DBMS

SQL features covered so far

- Query
- Modification
- Constraints
- Triggers
- Views
- Indexes